

## Question1:

## Python implementation of the algorithm:

```
[149] import numpy as np
import matplotlib.pyplot as plt

[129] class EpsilonGreedyBandit:
    def __init__(self, n_arms, epsilon):
        self.n_arms = n_arms
        self.epsilon = epsilon
        self.total_rewards = np.zeros(n_arms) # Keep track of total rewards for each arm
        self.action_counts = np.zeros(n_arms) # Keep track of number of times each arm is chosen
        self.average_rewards = np.zeros(n_arms) # Estimated average rewards for each arm

    def select_action(self):
        if np.random.random() < self.epsilon:
            # Explore: Select a random arm
            return np.random.choice(self.n_arms)
        else:
            # Exploit: Select the best arm
            return np.argmax(self.average_rewards)

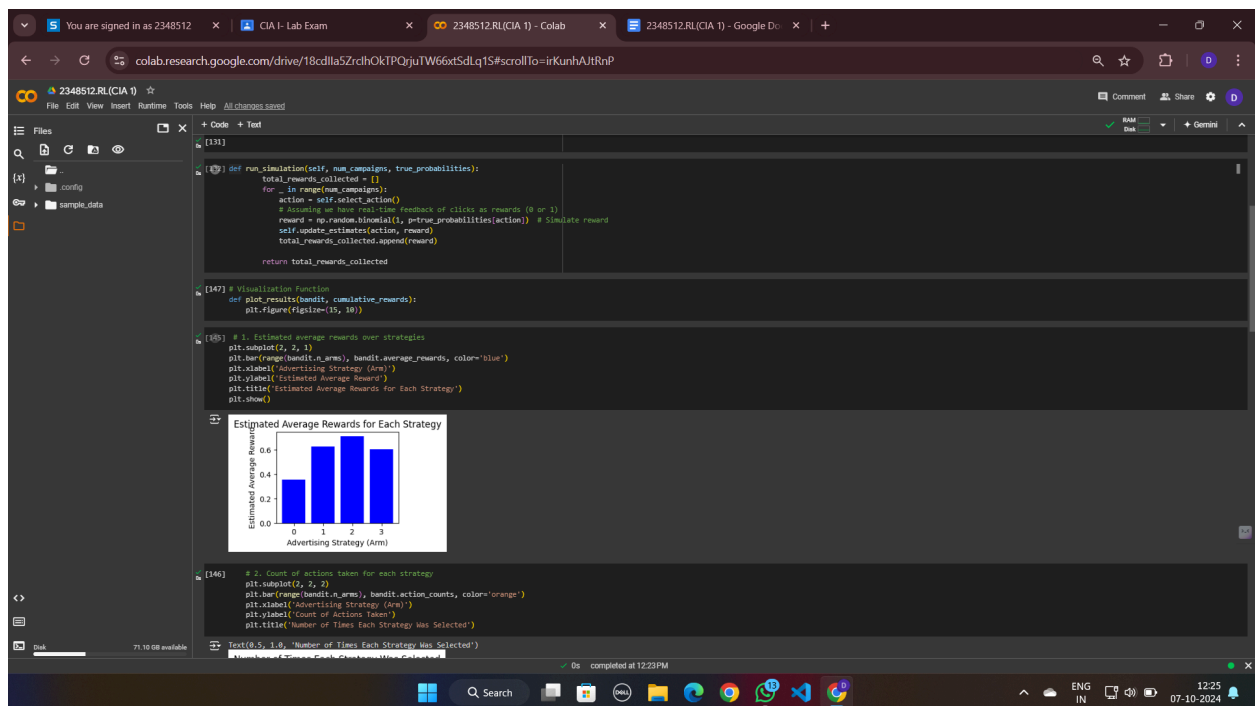
[131] def update_estimates(self, action, reward):
    # Update total rewards and counts for the action taken
    self.total_rewards[action] += reward
    self.action_counts[action] += 1
    # Update average rewards for the action
    self.average_rewards[action] = self.total_rewards[action] / self.action_counts[action]

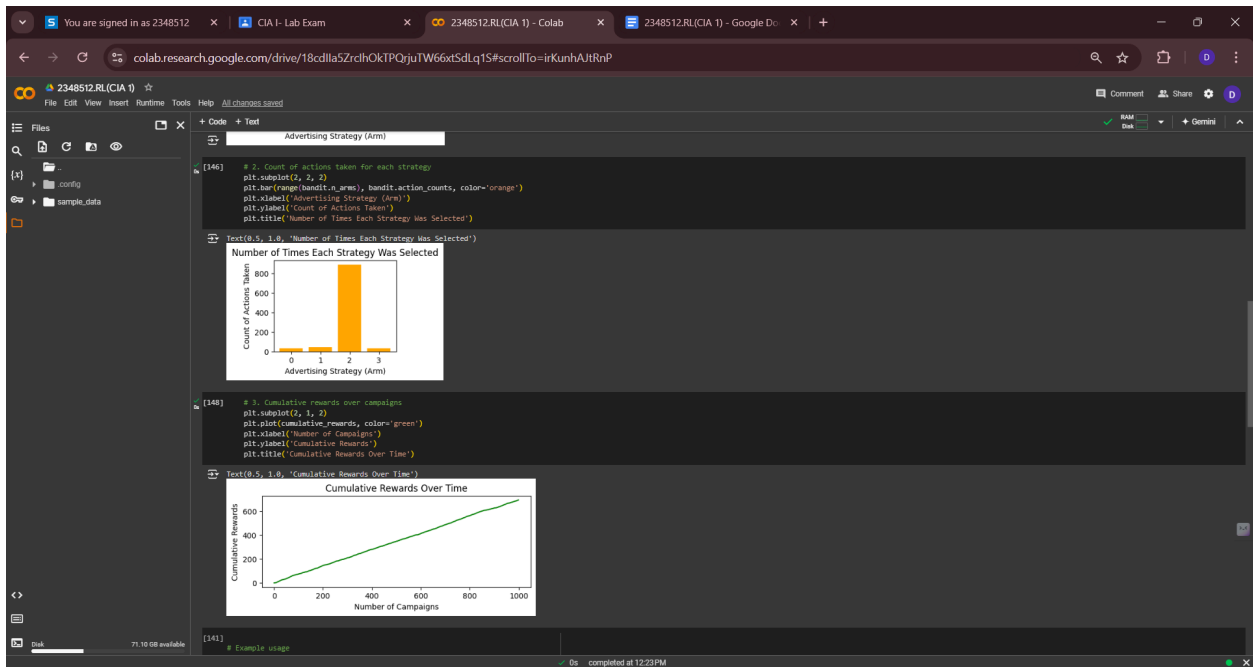
[132] def run_simulation(self, num_campaigns, true_probabilities):
    total_rewards_collected = []
    for _ in range(num_campaigns):
        action = self.select_action()
        # Assuming we have real-time feedback of clicks as rewards (0 or 1)
        reward = np.random.binomial(1, p=true_probabilities[action]) # Simulate reward
        self.update_estimates(action, reward)
        total_rewards_collected.append(reward)
    return total_rewards_collected

[147] # Visualization Function
def plot_results(bandit, cumulative_rewards):
    plt.figure(figsize=(15, 10))

[145] # 1. Estimated average rewards over strategies
plt.subplot(2, 2, 1)
plt.bar(range(bandit.n_arms), bandit.average_rewards, color='blue')
```

## Average Rewards over strategies:





Show how the algorithm helps to choose the best-performing advertising strategy over time:

